

Neural Guided Program Synthesis

AAAI'22 & JOSS'22


with N.Fijalkow, G.Lagarde, T.Matricon, K.Ellis, P.Ohlmann, A.Potta

Séminaire LINKS - 05/04/2024

Airport Screening


Population

Who to scan in order to catch terrorists?



Terrorists

Airport Screening



Who to scan in order to catch terrorists?

Assumption: prior knowledge
Say: someone in group A is 9
times more likely to be a terrorist

Airport Screening

We have a distribution D

We want to do sampling *with replacement* using another distribution, say D'

$\min_{D'} \mathbb{E}[\text{number of scans before finding a terrorist}]$

$Loss(D, D')$



Airport Screening

Naive: use D

Optimal: $\sqrt{D} = \frac{\sqrt{D(x)}}{\sum_y \sqrt{D(y)}}$

If group A is 9 times more likely to be a terrorist, it should be sampled 3 times more often.

Airport Screening

Simplest instance:

$$D : \quad p \cdot HEAD + (1 - p) \cdot TAIL$$

$$D' : \quad p' \cdot HEAD + (1 - p') \cdot TAIL$$


Airport Screening

Simplest instance:

$$D : p \cdot HEAD + (1 - p) \cdot TAIL$$

$$D' : p' \cdot HEAD + (1 - p') \cdot TAIL$$

$$Loss(D, D') = \frac{p}{p'} + \frac{1 - p}{1 - p'}$$



$$p' = \frac{\sqrt{p}}{\sqrt{p} + \sqrt{1 - p}}$$



Airport Screening

Simplest instance:

$$D : p \cdot HEAD + (1 - p) \cdot TAIL$$

$$D' : p' \cdot HEAD + (1 - p') \cdot TAIL$$


$$Loss(D, D') = \frac{p}{p'} + \frac{1 - p}{1 - p'}$$



$$p' = \frac{\sqrt{p}}{\sqrt{p} + \sqrt{1 - p}}$$


Program Synthesis?

An old dream: Church's Problem (1957)



Program Synthesis?

An old dream: Church's Problem (1957)



Logical formulas

Specification = ϕ a logical formula
A program P such that for all x , $\phi(x, P(x)) = \text{True}$

Natural language

« A program that removes odd elements and sort the rest »

A set of I/O examples

$[1, 5, 4, 2] \longrightarrow [2, 4]$
 $[6, 3, 0, 8] \longrightarrow [0, 6, 8]$

Today's setting! aka « *Inductive Program Synthesis* »

Program Synthesis?

An old dream: Church's Problem (1957)

Specification \longrightarrow **Program Synthesis** \longrightarrow Program

Logical formulas

Specification = ϕ a logical formula
A program P such that for all x , $\phi(x, P(x)) = \text{True}$

Natural language

« A program that removes odd elements and sort the rest »

A set of I/O examples

$[1, 5, 4, 2] \longrightarrow [2, 4]$
 $[6, 3, 0, 8] \longrightarrow [0, 6, 8]$

SORT; FILTER(EVEN)

Today's setting! aka « *Inductive Program Synthesis* »

Program Synthesis?

An old dream: Church's Problem (1957)

Specification \longrightarrow **Program Synthesis** \longrightarrow Program

Logical formulas

Specification = ϕ a logical formula
A program P such that for all x , $\phi(x, P(x)) = \text{True}$

Natural language

« A program that removes odd elements and sort the rest »

A set of I/O examples

$[1, 5, 4, 2] \longrightarrow [2, 4]$
 $[6, 3, 0, 8] \longrightarrow [0, 6, 8]$

Today's setting! aka « *Inductive Program Synthesis* »

Similarities with machine learning



Similarities with machine learning



Program Synthesis	Machine learning
Small number of examples	Large amount of data
Combinatorial/Symbolic search	Optimisation
Satisfies all specifications (reliable)	Minimizes a loss function (noisy)
Program (interpretable)	Model (hard to understand)

DeepCoder

Microsoft (Balog et al., 2017) — it manipulates list of integers

Program 4:

```
x ← [int]
y ← [int]
c ← SORT x
d ← SORT y
e ← REVERSE d
f ← ZIPWITH (*) d e
g ← SUM f
```

Input-output example:

Input:

```
[7 3 8 2 5],
[2 8 9 1 3]
```

Output:

```
79
```

Description:

Xavier and Yasmine are laying sticks to form non-overlapping rectangles on the ground. They both have fixed sets of pairs of sticks of certain lengths (represented as arrays x and y of numbers). Xavier only lays sticks parallel to the x axis, and Yasmine lays sticks only parallel to y axis. Compute the area their rectangles will cover at least.

TF-Coder

Google (Shi et al. 2019) — it manipulates tensors in Tensorflow





```
# Input-output example
inputs = {
    'scores': [[0.7, 0.2, 0.1],
               [0.4, 0.5, 0.1],
               [0.4, 0.4, 0.2],
               [0.3, 0.4, 0.3],
               [0.0, 0.0, 1.0]],
}
output = [[1, 0, 0],
          [0, 1, 0],
          [1, 0, 0],
          [0, 1, 0],
          [0, 0, 1]]
```

TF-Coder uses a combination of `tf.one_hot` and `tf.argmax` in a short solution to this problem:

```
tf.cast(tf.one_hot(tf.argmax(scores, axis=1), 3), tf.int32)
```


ARC Dataset

« *The Abstraction and Reasoning Corpus* », in « *On the measure of intelligence* » François Chollet, 2019



Neural-Guided Program Synthesis

Combination of **formal methods** and **machine learning**


reliable

efficient

Machine Learning Predictions

Symbolic Search

Evaluation



Neural-Guided Program Synthesis

Combination of **formal methods** and **machine learning**


reliable

efficient

Machine Learning Predictions

Symbolic Search

Evaluation



First practical instance: **DeepCoder, 2017**

Natural idea: patterns found in **examples** reveal which primitives are used in a solution program

[1, 5, 4, 2]	→	[2, 4]
[6, 3, 0, 8]	→	[0, 6, 8]

Likely primitives such that:

- SORT
- FILTER[EVEN]
- MULTIPLY[2]
- ...

Program representation

DSL: domain specific language

- Designed for a specific class of tasks
- List of primitives + associated types + semantics

DSL

```
sort : list(int) → list(int)
car : list(t) → t
cdr : list(t) → list(t)
map : (t1 → t2) →
      list(t1) → list(t2)
range : int → list(int)
+ : int → int → int
1 : int
⋮
```



The Problem

A Syntax-Guided Synthesis problem (SyGuS, in short) is specified with respect to a background theory \mathbb{T} , such as Linear-Integer-Arithmetic (LIA), that fixes the types of variables, operations on types, and their interpretation.

To synthesize a function f of a given type, the input consists of two constraints: **(1)** a *semantic constraint* given as a formula φ built from symbols in theory \mathbb{T} along with f , and **(2)** a *syntactic constraint* given as a (possibly infinite) set \mathcal{E} of expressions from \mathbb{T} specified by a context-free grammar. The computational problem then is to find an implementation for the function f , i.e. an expression $e \in \mathcal{E}$ such that the formula $\varphi[f \leftarrow e]$ is valid.

DSL \rightarrow Context-free Grammar

DSL

```
sort : list(int)  $\rightarrow$  list(int)
car : list(t)  $\rightarrow$  t
cdr : list(t)  $\rightarrow$  list(t)
map : (t1  $\rightarrow$  t2)  $\rightarrow$ 
      list(t1)  $\rightarrow$  list(t2)
range : int  $\rightarrow$  list(int)
+ : int  $\rightarrow$  int  $\rightarrow$  int
1 : int
       $\vdots$ 
```

Syntactic constraints

```
program type list(int)  $\rightarrow$  int
program depth  $\leq$  6
type size  $\leq$  10
type nesting  $\leq$  4
variable depth  $\geq$  3
no two consecutive sort
       $\vdots$ 
```

compilation

CFG

$S \rightarrow I, 0$

$I, d \rightarrow \text{one}$
 $I, d \rightarrow \text{car}_I(L[I], d+1)$
 $I, d \rightarrow \text{plus}(I, d+1; I, d+1)$
 \vdots
 $d=0, 1, 2, 3, 4, 5$

$L[I], d \rightarrow \text{range}(I, d+1)$
 $L[I], d \rightarrow \text{var} \quad \text{if } d \geq 3$
 $L[I], d \rightarrow \text{sort}(L[I], d+1)$
 $L[I], d \rightarrow \text{cdr}_I(L[I], d+1)$
 $L[I], d \rightarrow \text{car}_{L[I]}(L[L[[I]]], d+1)$
 $L[I], d \rightarrow \text{map}_{I,I}(I \text{ to } I, d+1;$
 $\quad \quad \quad L[I], d+1)$
 \vdots
 $d=0, 1, 2, 3, 4, 5$

$I \text{ to } I, d \rightarrow \text{car}_{I \text{ to } I}(L[I \text{ to } I], d+1)$
 $I \text{ to } I, d \rightarrow \text{plus}(I, d+1)$
 \vdots
 $d=0, 1, 2, 3, 4, 5$

From CFG to **P**CFG

CFG

$$S \rightarrow f(S, S) + g(T)$$

$$T \rightarrow a + b$$

From CFG to PCFG

CFG

$$S \rightarrow f(S, S) + g(T)$$

$$T \rightarrow a + b$$

PCFG

$$S \rightarrow \frac{1}{4} \cdot f(S, S) + \frac{3}{4} \cdot g(T)$$

$$T \rightarrow \frac{2}{3} \cdot a + \frac{1}{3} \cdot b$$

From CFG to PCFG

CFG

$$S \rightarrow f(S, S) + g(T)$$

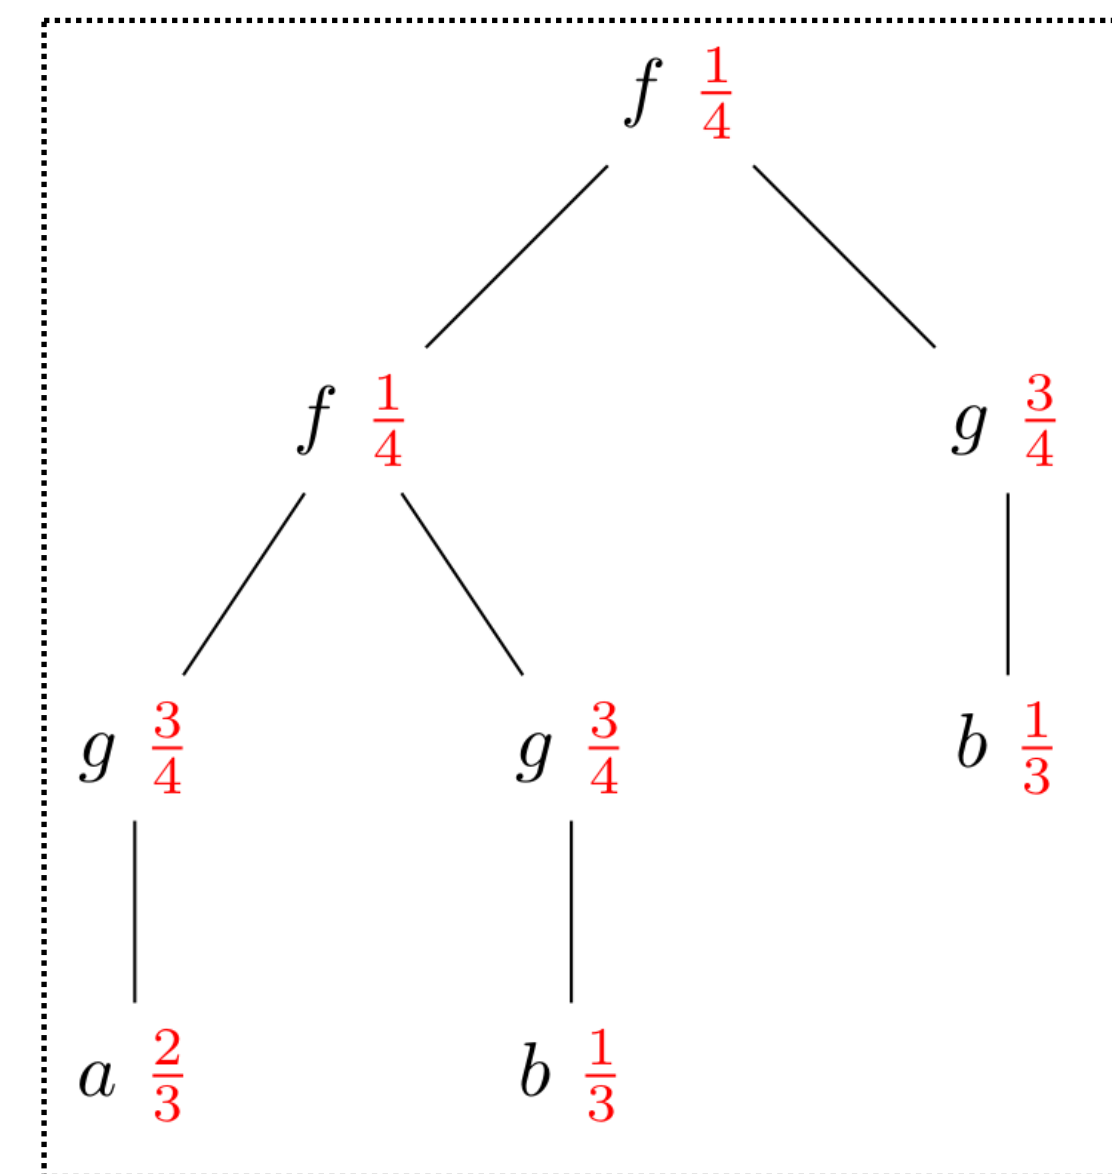
$$T \rightarrow a + b$$

PCFG

$$S \rightarrow \frac{1}{4} \cdot f(S, S) + \frac{3}{4} \cdot g(T)$$


$$T \rightarrow \frac{2}{3} \cdot a + \frac{1}{3} \cdot b$$

A PCFG induces a distribution D over **trees = programs**




$$\text{Probability} = \frac{1}{4} \cdot \frac{1}{4} \cdot \frac{3}{4} \cdot \frac{2}{3} \dots$$

Predictions: learning the PCFG



Predictions: learning the PCFG

Pr(SORT | FILTER) is high
Pr(SORT | REV) is low




A prediction
model

Probabilistic weights for the CFG

$$S \rightarrow \frac{1}{4} \cdot f(S, S) + \frac{3}{4} \cdot g(T)$$
$$T \rightarrow \frac{2}{3} \cdot a + \frac{1}{3} \cdot b$$

Predictions: learning the PCFG

$Pr(SORT | FILTER)$ is high
 $Pr(SORT | REV)$ is low



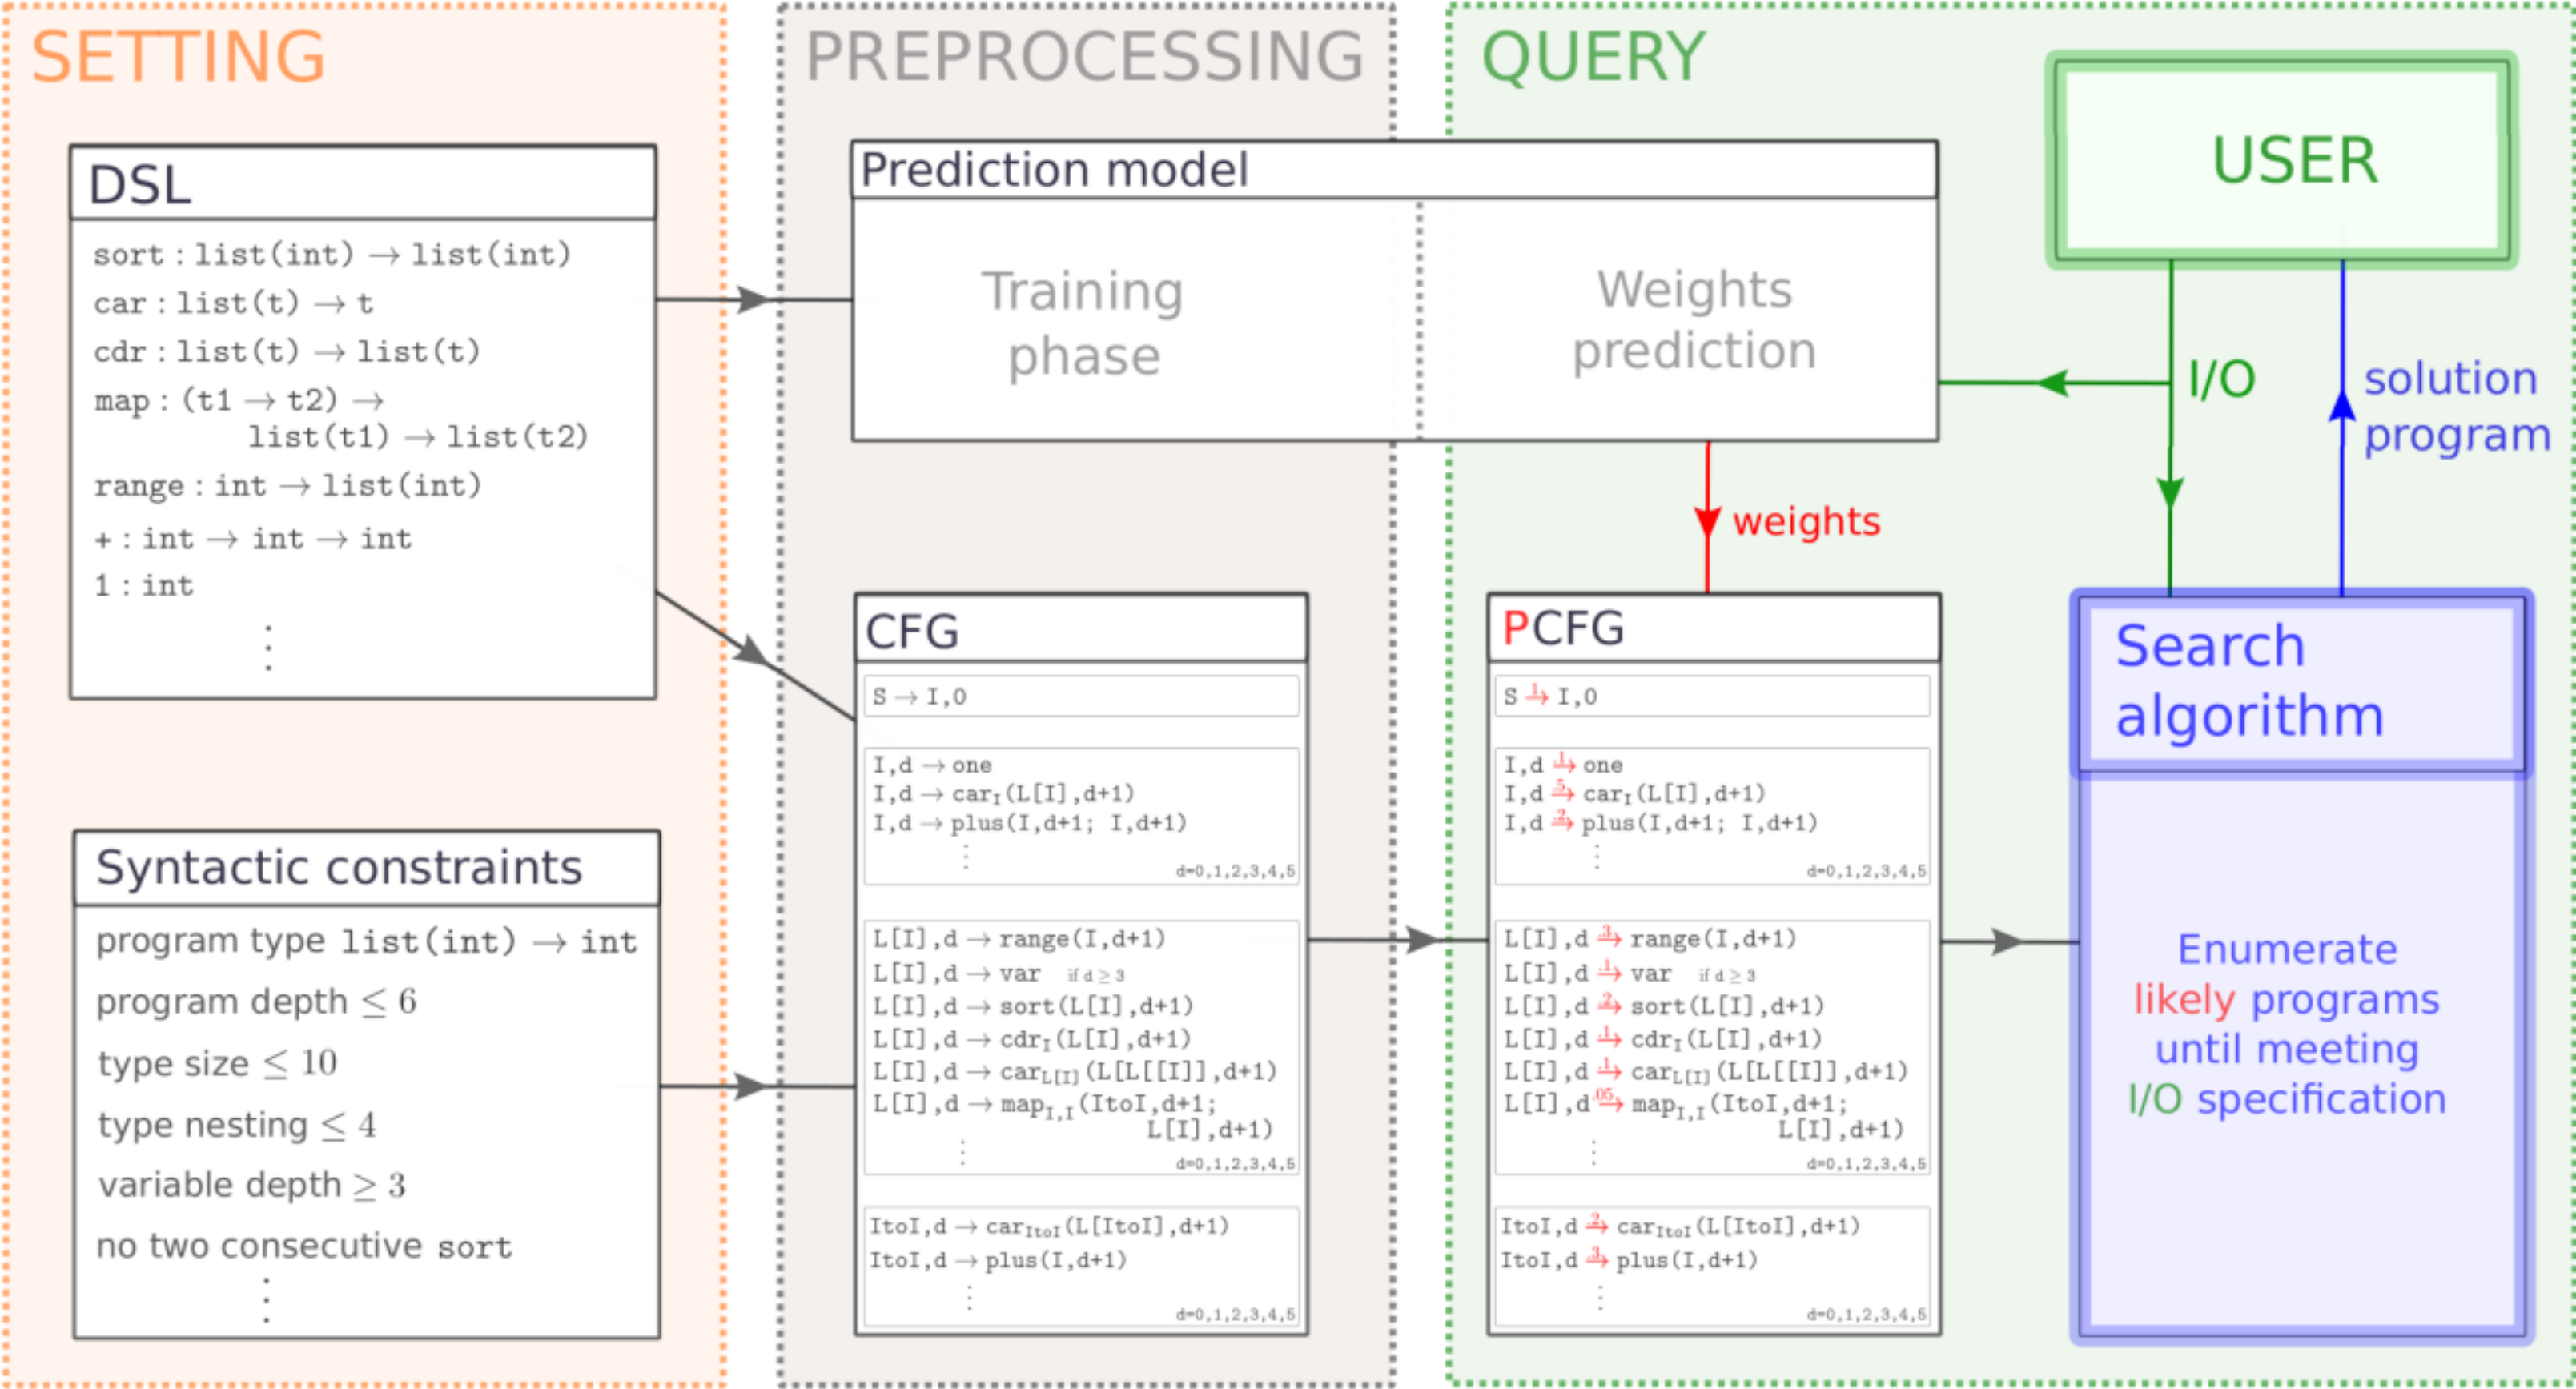
A prediction
model

Probabilistic weights for the CFG

$$S \rightarrow \frac{1}{4} \cdot f(S, S) + \frac{3}{4} \cdot g(T)$$
$$T \rightarrow \frac{2}{3} \cdot a + \frac{1}{3} \cdot b$$

The prediction model induces a **prior** distribution $D = Pr(program | examples)$

Summary so far



We are here!

How to use predictions for the search?

Toy example on the board

Technical contributions

Technical contributions

Heap Search.



Bottom-up enumeration strategy **preserving the order on programs**

- **Idea:** collections of heaps for storing partial programs
- **Guarantee:** i-th program in time $O(\log i)$

Grammar Splitting.

Divide the search on a parallel architecture

- **Idea:** partition derivations in a fair way



Technical contributions

Heap Search.

Bottom-up enumeration strategy **preserving the order on programs**



- **Idea:** collections of heaps for storing partial programs
- **Guarantee:** i -th program in time $O(\log i)$

Is it optimal?

Grammar Splitting.

Divide the search on a parallel architecture

- **Idea:** partition derivations in a fair way



Technical contributions

Heap Search.

Bottom-up enumeration strategy **preserving the order on programs**

- **Idea:** collections of heaps for storing partial programs
- **Guarantee:** i-th program in time $O(\log i)$

Is it optimal?

SQRT Sampling.

Optimal sampling strategy with **no memory**



Interesting for simplicity and parallel search

- **Idea:** sample programs according to $D' \propto \sqrt{D}$

Grammar Splitting.

Divide the search on a parallel architecture

- **Idea:** partition derivations in a fair way



Technical contributions

Heap Search.

Bottom-up enumeration strategy **preserving the order on programs**

- **Idea:** collections of heaps for storing partial programs
- **Guarantee:** i-th program in time $O(\log i)$

Is it optimal?

SQRT Sampling.

Optimal sampling strategy with **no memory**

Interesting for simplicity and parallel search



- **Idea:** sample programs according to $D' \propto \sqrt{D}$

How to sample? \rightarrow Construct a new PCFG for \sqrt{D}

Grammar Splitting.

Divide the search on a parallel architecture

- **Idea:** partition derivations in a fair way






Figure 7: Comparing all search algorithms on the DreamCoder reduced dataset with machine-learned PCFGs

Thanks!

- Generic pipeline to do **symbolic search** on grammars enhanced with **Machine Learning Predictions**
- Check out **DeepSynth** (open-source tool) — <https://deepsynth.labri.fr/>
- Is heap search optimal (from an enumeration complexity POV)?
- Better sampling if using memory?
- Use feedback to refine the search?

Example

$$G \begin{cases} S \rightarrow \frac{1}{2} f(S, T) \mid \frac{1}{2} a \\ T \rightarrow b \end{cases}$$

Natural candidate:

$$G' \begin{cases} S \rightarrow \frac{1}{\sqrt{2}} f(S, T) \mid \frac{1}{\sqrt{2}} a \\ T \rightarrow b \end{cases}$$

G' indeed computes \sqrt{D} but it is not a PCFG,
it needs renormalisation

for X non-terminal ($X = S, T$) define

$$[X] = \sum \{ D(t) : t \text{ generated from } X \}$$

$$\begin{cases} [S] = \frac{1}{\sqrt{2}} [S][T] + \frac{1}{\sqrt{2}} \\ [T] = 1 \end{cases} \Rightarrow \begin{cases} [S] = 1 + \sqrt{2} \\ [T] = 1 \end{cases}$$

The PCFG computing \sqrt{D} is

$$\sqrt{G} \begin{cases} S \rightarrow \frac{1}{1+\sqrt{2}} \left(\frac{1}{\sqrt{2}} f(S, T) \mid \frac{1}{\sqrt{2}} a \right) \\ T \rightarrow b \end{cases}$$